

Assignment 06

Section No. 1: Introduction

Briefly describe the classification problem and general data preprocessing. Note that some data preprocessing steps maybe specific to a particular algorithm. Report those steps under each algorithm section.

For this assignment, we are viewing a dataset with digits from 0 to 9 of all handwritten images. We will use both the decision tree and naïve bayes to see which model gives us the best prediction.

```
#import the following packages
```

```
library(dplyr)
```

```
library(caret)
```

```
library(rpart)
```

```
library(e1071)
```

```
library(stringr)
```

```
#import data
```

```
digittrain <- read.csv("/Volumes/STORE N GO/00 - Graduate School/00 - SYR In Session/00 - Fall 2019/IST  
707 - Data Analytics Wed9PM/Deliverables/HWs/HW6/digittrain.csv")
```

```
digittest <- read.csv("/Volumes/STORE N GO/00 - Graduate School/00 - SYR In Session/00 - Fall 2019/IST  
707 - Data Analytics Wed9PM/Deliverables/HWs/HW6/digittest.csv")
```

```
dim(digittrain)
```

We observe that the training dataset has 42,000 observations and 785 variables.

```
str(digittrain[, 1:10])
```

```
> str(digittrain[, 1:10])
```

```
'data.frame': 42000 obs. of 10 variables:  
 $ label : Factor w/ 10 levels "0","1","2","3",...: 2 1 2 5 1 1 8 4 6 4 ...  
 $ pixel0: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel1: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel2: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel3: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel4: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel5: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel6: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel7: int 0 0 0 0 0 0 0 0 0 0 ...  
 $ pixel8: int 0 0 0 0 0 0 0 0 0 0 ...
```

```
summary(digittrain[, 1:10])
```

```
dim(digittest)
```

We observe that the test dataset has 28,000 observations and 784 variables.

```
str(digittest[, 1:10])  
summary(digittest[, 1:10])
```

The dataset is 0 through 9 handwritten images & 784-pixel variables. Our labels are "Pixels" and our "data type are all integers

We will be working with a sample of the data (due to the size) for building our model. As described in the "Task Description", we will be sampling 10% of our training and testing data to train and test to be used for our classifier."

```
#we set the seed for randomness  
set.seed(150)
```

```
# selection of 10% of the training & test data  
splitrain <- sample(nrow(digittrain), nrow(digittrain) * .1)  
splitest <- sample(nrow(digittest), nrow(digittest) * .1)
```

```
# row selection to build our set  
subtrain <- digittrain[splitrain, ]  
subtest <- digittest[splitest, ]
```

```
#specify that the row names are "NULL", that there are none  
row.names(digittrain) <- NULL  
row.names(digittest) <- NULL
```

We will also attempt to view the heading of our datasets.

Section No. 2: Decision Tree

Build a decision tree model. Tune the parameters, such as the pruning options, and report the 3-fold CV accuracy.

Decision tree model

```
digittreetrain <- rpart(label ~ ., data = subtrain, method = 'class',
  control = rpart.control(cp = 0), minsplit = 100, maxdepth = 10)
```

```
> digittreetrain
n= 4200
```

```
node), split, n, loss, yval, (yprob)
* denotes terminal node
```

```
1) root 4200 3745 1 (0.096 0.11 0.1 0.1 0.1 0.095 0.095 0.1 0.098 0.095)
  2) pixel409< 0.5 1378 970 1 (0.24 0.3 0.078 0.061 0.022 0.13 0.052 0.028 0.08 0.015)
    4) pixel462< 3 512 184 0 (0.64 0.0059 0.059 0.066 0.0039 0.09 0.053 0.043 0.021 0.018)
      8) pixel455>=9.5 319 40 0 (0.87 0 0.019 0.0094 0.0063 0.025 0.034 0.0063 0.0063 0.019)
        16) pixel351< 161.5 295 23 0 (0.92 0 0.017 0.0034 0.0068 0.014 0.01 0.0068 0 0.02)
          32) pixel517< 242.5 280 13 0 (0.95 0 0.0036 0.0036 0.0036 0.0071 0.0036 0.0071 0 0.018)
            64) pixel628>=18.5 250 3 0 (0.99 0 0.004 0.004 0 0.004 0 0 0) *
              65) pixel628< 18.5 30 10 0 (0.67 0 0 0.033 0.033 0.033 0.067 0 0.17)
                130) pixel442>=25 23 4 0 (0.83 0 0 0 0.043 0.043 0 0.087) *
                  131) pixel442< 25 7 4 9 (0.14 0 0 0.14 0.14 0.14 0 0.43) *
                    33) pixel517>=242.5 15 10 0 (0.33 0 0.27 0.067 0.13 0.13 0 0.067) *
                      17) pixel351>=161.5 24 16 6 (0.29 0 0.042 0.083 0 0.17 0.33 0 0.083 0)
                        34) pixel515< 55 16 9 0 (0.44 0 0.062 0.12 0 0.25 0 0.12 0) *
                          35) pixel515>=55 8 0 6 (0 0 0 0 0 1 0 0 0) *
                            9) pixel455< 9.5 193 144 0 (0.25 0.016 0.12 0.16 0 0.2 0.083 0.1 0.047 0.016)
                              18) pixel324< 13.5 88 42 0 (0.52 0.023 0.11 0.034 0 0.14 0.1 0.068 0 0)
                                36) pixel517< 16.5 68 22 0 (0.68 0.015 0.029 0.044 0 0.1 0.044 0.088 0 0)
                                  72) pixel405< 38.5 56 10 0 (0.82 0 0.018 0 0.018 0.036 0.11 0 0) *
                                    73) pixel405>=38.5 12 6 5 (0 0.083 0.083 0.25 0 0.5 0.083 0 0) *
                                      37) pixel517>=16.5 20 12 2 (0 0.05 0.4 0 0.25 0.3 0 0 0)
                                        74) pixel239< 11 9 1 2 (0 0 0.89 0 0.11 0 0 0) *
                                          75) pixel239>=11 11 5 6 (0 0.091 0 0 0.36 0.55 0 0 0) *
                                            19) pixel324>=13.5 105 77 3 (0.029 0.0095 0.13 0.27 0 0.25 0.067 0.13 0.086 0.029)
                                              38) pixel514>=49 36 22 2 (0.028 0.028 0.39 0.028 0 0.083 0.19 0 0.25 0)
                                                76) pixel436< 148 26 12 2 (0.038 0.038 0.54 0.038 0 0.12 0.23 0 0 0)
                                                  152) pixel355< 6 15 1 2 (0 0.067 0.93 0 0 0 0 0 0) *
                                                    153) pixel355>=6 11 5 6 (0.091 0 0 0.091 0 0.27 0.55 0 0 0) *
                                                      77) pixel436>=148 10 1 8 (0 0 0 0 0.1 0.9 0) *
                                                        39) pixel514< 49 69 42 3 (0.029 0 0.39 0 0.33 0 0.2 0 0.043)
                                                          78) pixel287< 175 52 26 3 (0.038 0 0 0.5 0 0.4 0 0.019 0 0.038)
                                                            156) pixel262< 11.5 30 7 3 (0.033 0 0.77 0 0.13 0 0.033 0 0.033)
                                                              312) pixel152>=7 20 0 3 (0 0 1 0 0 0 0 0) *
                                                                313) pixel152< 7 10 6 5 (0.1 0 0.3 0 0.4 0 0.1 0 0.1) *
                                                                  157) pixel262>=11.5 22 5 5 (0.045 0 0.14 0 0.77 0 0 0.045) *
                                                                    79) pixel287>=175 17 4 7 (0 0 0.059 0 0.12 0 0.76 0 0.059) *
                                                                      5) pixel462>=3 866 461 1 (0.0046 0.47 0.089 0.058 0.033 0.15 0.051 0.02 0.11 0.014)
                                                                        10) pixel347< 1.5 570 180 1 (0 0.68 0.11 0.044 0.023 0.046 0.014 0.028 0.039 0.0088)
                                                                          20) pixel551< 6 476 90 1 (0 0.81 0.036 0.025 0.0084 0.042 0.0063 0.029 0.038 0.0042)
                                                                            40) pixel350>=39.5 430 49 1 (0 0.89 0.012 0.023 0.0047 0.012 0.007 0.016 0.04 0)
                                                                              80) pixel300< 75 412 32 1 (0 0.92 0.012 0.015 0.0049 0.0097 0.0073 0.017 0.012 0)
                                                                                160) pixel494< 48 402 22 1 (0 0.95 0.005 0.012 0.0025 0.01 0.0075 0.005 0.012 0)
                                                                                  320) pixel206< 48.5 392 15 1 (0 0.96 0 0.0026 0 0.01 0.0077 0.0051 0.013 0) *
                                                                                    321) pixel206>=48.5 10 6 3 (0 0.3 0.2 0.4 0.1 0 0 0 0) *
                                                                                      161) pixel494>=48 10 5 7 (0 0.3 0.1 0.1 0.1 0.5 0 0) *
                                                                                        81) pixel300>=75 18 6 8 (0 0.056 0 0.22 0 0.056 0 0 0.67 0) *
                                                                                          41) pixel350< 39.5 46 31 5 (0 0.11 0.26 0.043 0.043 0.33 0 0.15 0.022 0.043)
```

Testing the accuracy of the tree on the training set

```
trainacc <- data.frame(predict(digittreetrain, subtrain))
```

#choose max likelihood

```
trainacc <- as.data.frame(names(trainacc[apply(trainacc, 1, which.max)]))
```

```
colnames(trainacc) <- 'prediction'
```

```
trainacc$number <- substr(trainacc$prediction, 2, 2)
```

```
trainacc <- subtrain %>% bind_cols(trainacc) %>% select(label, number) %>% mutate(label =
as.factor(label), number = as.factor(round(as.numeric(number), 0)))
```

Confusion Matrix

Now let's build the Confusion matrix so we can examine the accuracy percentage

```
confusionMatrix(trainacc$label, trainacc$number)
```

#our training set gives us an 86% (approximate) accuracy.

```
> confusionMatrix(trainacc$label, trainacc$number)
Confusion Matrix and Statistics
```

	Reference									
Prediction	0	1	2	3	4	5	6	7	8	9
0	373	0	3	8	6	4	2	1	5	3
1	0	415	4	9	9	3	6	7	2	0
2	12	4	360	12	12	4	5	5	3	7
3	4	1	7	338	10	26	5	6	18	12
4	6	2	5	11	386	3	3	2	6	17
5	10	6	3	22	7	318	11	6	5	12
6	6	3	4	3	13	5	362	0	2	1
7	7	3	10	3	21	4	1	372	4	13
8	5	6	9	8	14	8	5	1	346	11
9	6	0	6	5	26	8	2	7	4	334

Overall Statistics

```
Accuracy : 0.8581
95% CI : (0.8472, 0.8685)
No Information Rate : 0.12
P-Value [Acc > NIR] : < 2.2e-16
```

```
Kappa : 0.8423
```

```
Mcnemar's Test P-Value : NA
```

Statistics by Class:

	Class: 0	Class: 1	Class: 2	Class: 3	Class: 4	Class: 5	Class: 6
Sensitivity	0.86946	0.94318	0.87591	0.80668	0.7659	0.83029	0.90050
Specificity	0.99151	0.98936	0.98311	0.97646	0.9851	0.97852	0.99026
Pos Pred Value	0.92099	0.91209	0.84906	0.79157	0.8753	0.79500	0.90727
Neg Pred Value	0.98524	0.99332	0.98649	0.97853	0.9686	0.98289	0.98948
Prevalence	0.10214	0.10476	0.09786	0.09976	0.1200	0.09119	0.09571
Detection Rate	0.08881	0.09881	0.08571	0.08048	0.0919	0.07571	0.08619
Detection Prevalence	0.09643	0.10833	0.10095	0.10167	0.1050	0.09524	0.09500
Balanced Accuracy	0.93049	0.96627	0.92951	0.89157	0.8755	0.90440	0.94538

	Class: 7	Class: 8	Class: 9
Sensitivity	0.91400	0.87595	0.81463
Specificity	0.98260	0.98239	0.98311
Pos Pred Value	0.84932	0.83777	0.83920
Neg Pred Value	0.99070	0.98706	0.98001
Prevalence	0.09690	0.09405	0.09762
Detection Rate	0.08857	0.08238	0.07952
Detection Prevalence	0.10429	0.09833	0.09476
Balanced Accuracy	0.94830	0.92917	0.89887

now we run the same prediction on the test data

```
testacc <- data.frame(predict(digittreetrain, subtest))
testacc <- as.data.frame(names(testacc[apply(testacc, 1, which.max)]))
colnames(testacc) <- 'prediction'
testacc$number <- substr(testacc$prediction, 2, 2)
```


Seems about even but we cannot test the results unless we run the decision tree on the entire test set and submit it to Kaggle.

```
finaltest <- data.frame(predict(digittreetrain, digittest))
finaltest <- as.data.frame(names(finaltest[apply(finaltest, 1, which.max)]))
colnames(finaltest) <- 'ImageId'
finaltest$Label <- substr(finaltest$ImageId, 2, 2)
finaltest$ImageId <- 1:nrow(finaltest)
```

#Now we can export the model (file) and view Kaggle results

```
write.csv(finaltest, file="/Volumes/STORE N GO/00 - Graduate School/00 - SYR In Session/00 - Fall 2019/IST
707 - Data Analytics Wed9PM/Deliverables/HWs/HW6/FEsuruoso - Digit Decision Tree Test.csv")
```

Sample of our Kaggle submission. We recorded an 81% Kaggle score (approximate).



✓ **FEsuruoso - Digit Decision Tree Test.csv (397.8 kB)**

File Format

Your submission should be in CSV format. You can upload this in a zip/gz/rar/7z archive, if you prefer.

Number of Predictions

We expect the solution file to have 28000 prediction rows. This file should have a header row. Please see sample submission file on the [data page](#).

↶ ↷ ⌂ B I 🔗 ” <> ☰ ☱ ☲ ☳ ☴ ☵ ☶ ☷ 🖼️ 😊

I processed my data and split it into a test and train data set. Due to the large size of the dataset, I sampled 10% of the training set to use for training and testing data to train and test to be used for our classifier.

Section No. 3: Naïve Bayes

Instruction: Build a naïve Bayes model. Tune the parameters, such as the discretization options, to compare results.

```
#build classifier
digittrainnb <- naiveBayes(as.factor(label) ~ ., data = subtrain)

nbtrainacc <- predict(digittrainnb, subtrain, type = 'class')
confusionMatrix(nbtrainacc, as.factor(subtrain$label))
```

```
> confusionMatrix(nbtrainacc, as.factor(subtrain$label))
Confusion Matrix and Statistics
```

	(No matches)	0	1	2	3	4	5	6	7	8	9
Prediction 0	391	0	83	71	32	104	11	8	10	15	
1	1	443	22	27	13	32	18	24	99	35	
2	0	0	75	2	0	1	0	1	3	0	
3	0	0	54	180	1	12	0	1	7	0	
4	0	0	4	1	75	0	2	1	2	1	
5	0	0	0	0	0	14	0	0	0	0	
6	7	5	99	13	65	15	360	2	3	9	
7	0	0	1	6	16	5	0	367	0	73	
8	6	7	85	121	166	207	8	23	284	59	
9	0	0	1	6	73	10	0	11	5	206	

```
Overall Statistics

      Accuracy : 0.5702
    95% CI : (0.5551, 0.5853)
 No Information Rate : 0.1083
 P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 0.5225

McNemar's Test P-Value : NA

Statistics by Class:
```

	Class: 0	Class: 1	Class: 2	Class: 3	Class: 4	Class: 5	Class: 6	Class: 7	Class: 8	Class: 9
Sensitivity	0.96543	0.9736	0.17689	0.42155	0.17007	0.035000	0.90226	0.83790	0.68765	0.51759
Specificity	0.91199	0.9276	0.99815	0.98012	0.99707	1.000000	0.94265	0.97315	0.81991	0.97212
Pos Pred Value	0.53931	0.6204	0.91463	0.70588	0.87209	1.000000	0.62284	0.78419	0.29400	0.66026
Neg Pred Value	0.99597	0.9966	0.91525	0.93739	0.91104	0.907788	0.98923	0.98098	0.96011	0.95062
Prevalence	0.09643	0.1083	0.10095	0.10167	0.10500	0.095238	0.09500	0.10429	0.09833	0.09476
Detection Rate	0.09310	0.1055	0.01786	0.04286	0.01786	0.003333	0.08571	0.08738	0.06762	0.04905
Detection Prevalence	0.17262	0.1700	0.01952	0.06071	0.02048	0.003333	0.13762	0.11143	0.23000	0.07429
Balanced Accuracy	0.93871	0.9506	0.58752	0.70083	0.58357	0.517500	0.92245	0.90553	0.75378	0.74485

Here we only see a 57% (approximate) test accuracy. This is not as good as the previous model.

```
nbtestacc <- predict(digittrainnb, digittest, type = 'class')
nbtestacc <- data.frame(nbtestacc)
colnames(nbtestacc)[1] <- 'Label'
nbtestacc$Imageld <- 1:nrow(nbtestacc)
nbtestacc <- nbtestacc %>% select(Imageld, Label)
```

```
# We will now submit our Naive Bayes algorithm to kaggle. First we need to export to csv
write.csv(nbtestacc, file="/Volumes/STORE N GO/00 - Graduate School/00 - SYR In Session/00 - Fall
2019/IST 707 - Data Analytics Wed9PM/Deliverables/HWs/HW6/FEsuruoso - Naive Bayes Classifier.csv",
row.names = FALSE)
```

Section No. 4: Algorithm Performance Comparison

Report the test accuracy for the naïve Bayes and decision tree models. Discuss whether overfitting occurs in these models.

I was not pleased with the results obtained by the Naïve Bayes algorithm. I received a 57% in my confusion matrix and a 57% on my Kaggle submission. I did not perform any pruning measures like I did with the previous model, so this might have a lot to do with it. There also weren't any additional parameters set.

One of the most significant measures of how good our models did was the Confusion Matrix accuracy. While we had around 80% for the first model, our second model only gave us around 57%. We see that the first model also properly identified our models better. It is also worthy to note that had we used a bigger portion of our data (we only used 10%) and tuned better, we might have had a better model.

Section No. 5: Decision Tree & Naïve Bayes Kaggle Test Result

Digit Recognizer | Kaggle

</

As shown above, I scored approximately 75% on my decision tree submission and 57% on my Naïve Bayes submission. Both scores could be improved greatly by the following:

- Increasing the sample size in our models from 10% to a larger sample size (maybe 20% or 40%)
- Additional tuning that addresses model fitting
- Address specific model issues (i.e smaller dataset for Naïve Bayes and larger dataset for decision tree)